



SmashClean: A Hardware level mitigation to stack smashing attacks in OpenRISC

Manaar Alam, Debapriya Basu Roy, Sarani Bhattacharya, Vidya Govindan

Rajat Subhra Chakraborty and Debdeep Mukhopadhyay

alam.manaar@gmail.com, vidya.mazhur@gmail.com, sarani.bhattacharya@cse.iitkgp.ernet.in, deb.basu.roy@cse.iitkgp.ernet.in, rschakraborty@cse.iitkgp.ernet.in, debdeep@cse.iitkgp.ernet.in

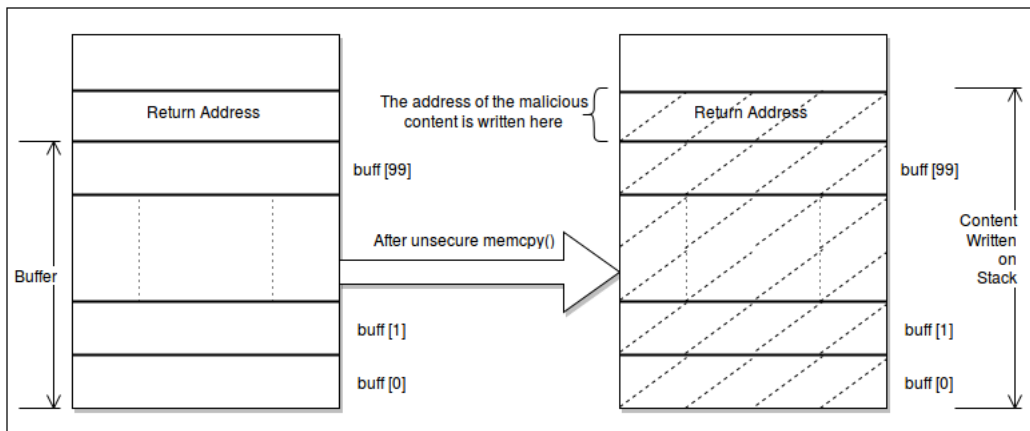
Secured Embedded Architecture Lab, Department of Computer Science and Engineering, Indian Institute of Technology Kharagpur, India

INTRODUCTION

- **Security threats to embedded systems**
 - Hardware and Software vulnerabilities
 - Performance-efficient languages such as C and C++ widely used for embedded applications
 - Vulnerable to memory corruption due to lack of secure management
- **Buffer Overflow:** Trigger malicious code execution by overwriting correct memory content
 - Software level countermeasures may be easily bypassed
 - Need hardware level countermeasures, e.g. hardware-based protection of the function return address
 - target platform for existing architectures different from the OpenRISC ISA processor

ATTACKING CONTROL FLOW

Return Address Modification



```
int func(char* user, int len) {
    char buff[100];
    memcpy(buff, user, len); //Vulnerability
}
```

Format String Vulnerability

```
int n;
printf("%12c%n", 'A', &n);
```

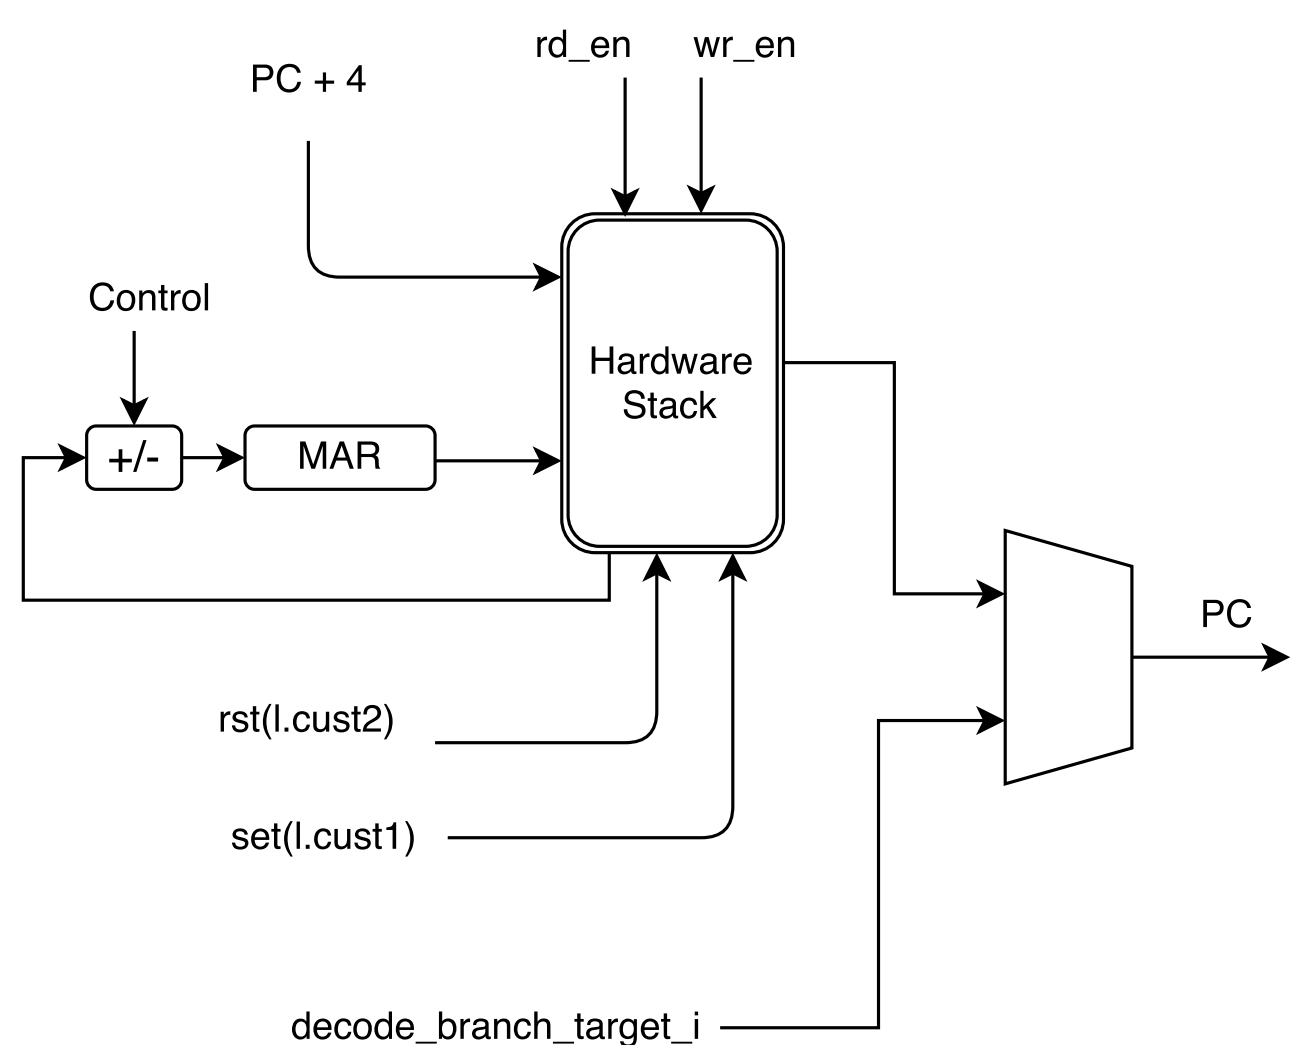
```
int func(char* user) {
    printf(user); //Vulnerability
}
```

Example: Assembly Code for Stack.c

```
vuln:
.LFB1:
.cfi_startproc
    .
    .
    .
    l.ori r1,r2,0 # deallocate frame
    l.lwz r2,-8(r1) # SI load
    l.lwz r9,-4(r1) # SI load
    l.jr r9 # return_internal
    l.nop # nop delay slot
.cfi_endproc
```

- If the address provided by a malicious user causes buffer overflow to modify r_9 then the control flow gets transferred to the malicious code

OUR HARDWARE STACK



OUR OBJECTIVE

Hardware-Based Mitigation of Memory Corruption and Ensuring Control Flow Integrity for the OpenRISC ISA Processor

OUR CONTRIBUTIONS

- Prevention of all forms of memory corruption and buffer overflow attacks on OpenRISC architecture
- Combination of compiler and hardware modification
- Introduction of new instructions via hardware modification for compiler to detect and prevent memory corruption via buffer overflow

PROTECT CONTROL FLOW

- Implementation of a hardware stack which stores the function return address for each of the function
- Prevention using hardware stack:
 - Whenever it encounters a `l.jal` or `l.jalr` instruction, it pushes the next program counter value to the stack
 - Alternatively if it encounters `l.jr` instruction with register r_9 as parameter, it pops its top value and passes that as the return address
 - Custom instruction `l.cust1`, when enabled, ensures that the return address of the functions are read from the hardware stack.
 - Custom instruction `l.cust2` disables the hardware stack.

PREVENT MEM. CORRUPTION


- We introduced hardware enforced secure `memcpy`
- This protection prevents buffer overflow by hardware induced bound check and prevents any memory corruption due to buffer overflow.

Example: Assembly Code for Priv.c

```
vuln:
l.sw -40(r2),r3 # SI store
.
.
.
l.sw -36(r2),r3 # SI store
.
.
.
l.nop # nop delay slot
l.lwz r4,-44(r2) # SI load
l.addi r3,r2,-32 # addsi3
l.ori r5,r4,0 # move reg to reg
l.lwz r4,-40(r2) # SI load
l.jal memcpy # call_value_internal
l.nop # nop delay slot
```

- The first instruction (`l.addi r3,r2,-32`) transfers the starting address of the buffer ($r_2 - 32$) to r_3 . The address of the latest new variable in this case is $r_2 - 16$. Subtracting this two will give us buffer size which in this case is 16.
- The next instruction `l.ori` transfers the function argument count to r_5 which denotes the number of memory locations to be updated by `memcpy`.
- Now, we will check whether the instruction `l.ori r5,r4,0` returns the count value greater than the buffer size or not.

Secure memcpy

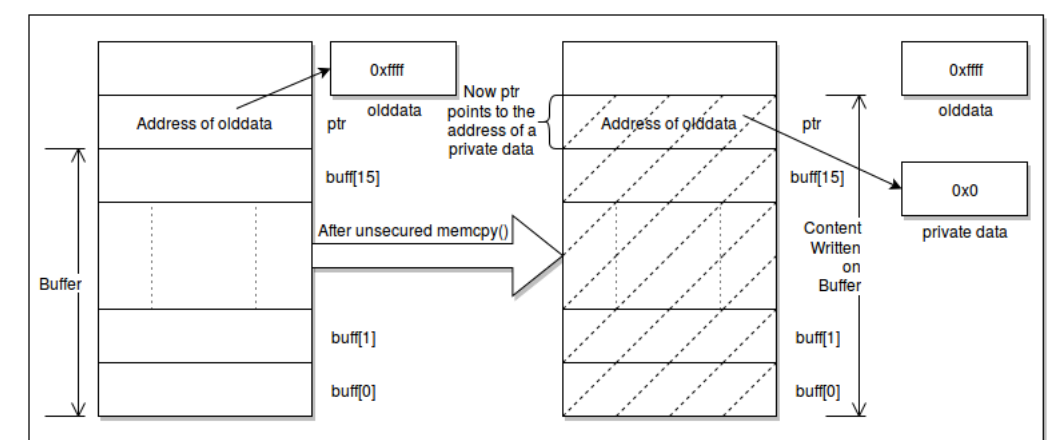


OVERVIEW

- The root cause of buffer overflow threat:
 - **memcpy does not impose any bound-checking during memory update**
- Our countermeasure approach:
 - Introduction of new instructions that keep track of the buffer size
 - * Ensures number of memory locations upgraded by `memcpy` is less than or equal to buffer size
 - Storing the return addresses in the hardware stack
 - * Prevention of return address modification
- Our results:
 - Prevented `stack.c`, `ptr.c` and `priv.c` on Linux platform using our new instruction
 - Protection against return address modification by `format.c` using stack-based tracking of return addresses

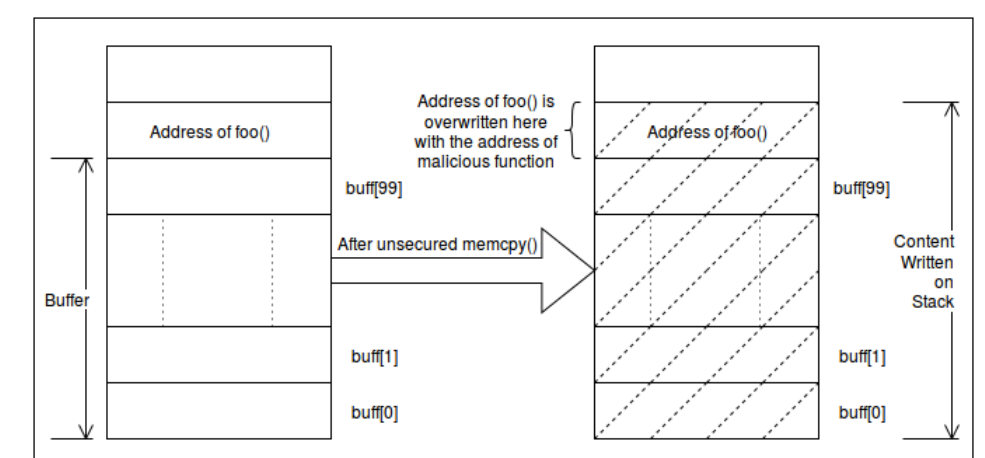
MEMORY CORRUPTION

Data Pointer Modification



```
int func(char* user, int len) {
    int *ptr;
    int newdata = 0xaaaa;
    char buff[16];
    int olddata = 0xffff;
    ptr = &olddata;
    memcpy(buff, user, len); //Vulnerability
    *ptr = newdata;
}
```

Function Pointer Modification



```
int func(char* user, int len) {
    void (*fptr)(char *);
    char buff[100];
    fptr = &foo; //Address of intended function
    memcpy(buff, user, len); //Vulnerability
    fptr(user);
}
```

NEW INSTRUCTIONS

- **l.cust3** This instruction will be inserted by the compiler just before `memcpy` function is declared in C code to protect buffer overflow. This instruction sets a specific flag inside the processor and observes the occurrence of `l.addi` and `l.ori` which are required for computation of buffer size. If the buffer size is less than the argument count a `smash_detect` flag is set and the value of the count argument is updated with the buffer size. Thus this instruction ensures both detection and prevention of buffer overflow.
- **l.cust4** This instruction resets the `smash_detect` flag.
- **l.cust5** This instruction induces a lock on latest variable address location to preserve it from intermediate function calls. This can be alternatively achieved by maintaining a hardware stack for latest variable locations for each function call.
- **l.cust6** This instruction removes the aforementioned lock.